

CSS – CIA 2 NOTES

UNIT 2 :

Policy Requirements

1. Users use only existing production programs.
2. Programmers develop/test on non-production systems.
3. Production data access via a special process.
4. Deployment to production follows a controlled process.
5. Deployment must be audited.
6. Managers/auditors get system state & logs access.

Integrity Policies : ensure data remains accurate and unaltered by unauthorized users.

1. Biba Integrity Model (Prevents data corruption)

- **Simple Integrity Property: No read down** (Higher-level subjects can't read lower-level data).
- **Star (*) Integrity Property: No write up** (Lower-level subjects can't modify higher-level data).

Example: A manager (high-level) can't read an intern's (low-level) draft, and an intern can't edit an executive's report.

2. Low-Water-Mark Model (Flexible version of Biba)

- Allows read down but lowers subject's integrity level if they do so.
- Prevents write up (similar to Biba).

Example: If a senior engineer reads an intern's code, their integrity level drops, restricting future actions.

3. Ring Policy Model (Less restrictive)

- Allows read down (unlike Biba).
- Subjects can only write at their own level.

Example: A software developer can read an intern's documentation but can only modify files at their assigned level.

4. Lipner's Integrity Matrix (Combines Biba + Bell-LaPadula)

- **No Read Down** (Prevents higher-level users from reading lower-integrity data).
- **No Write Up** (Prevents low-level users from corrupting high-integrity data).

Example: A financial officer can access payroll data but can't modify security system configurations.

Two Security Levels:

- **Audit Manager** → High integrity, controls security audits.
- **System Low** → Lower integrity, general operations.

Five Categories:

1. Development (D) – Application development environment.
2. Production Code (PC) – **Finalized software ready for deployment.**
3. Production Data (PD) – Real-time operational data.
4. System Development (SD) – **System-level software development.**
5. **Software Tools (T)** – Utilities used for development and maintenance.

Security levels for subjects

Users	Clearance
Ordinary Users	(SL,{PC,PD})
Application developers	(SL,{D,T})
System programmers	(SL,{SD,T})
System managers and auditors	(AM{D,PC,PD,SD,T})
System controllers	(SL,{D,PC,PD,SD,T}) downgrade privilege

Security levels for objects

Objects : Development code, production code, production data, software tools, system programs, system programs in modification, system and application logs

Security Requirements of the Model

- **Prevents Unauthorized Modification:** Only trusted users can modify high-integrity objects.
- **Separation of Duties:** Developers cannot directly modify production systems.
- **Auditability:** The Audit Manager monitors all integrity violations.

5. Lipner's Full Integrity Matrix (Combines Biba's Integrity Model for integrity & Bell-LaPadula Model for confidentiality)

- **Ensure Data Integrity:** Prevent unauthorized or accidental modifications.
- **Separate Development & Production:** Developers can't alter production data.
- **Maintain Auditability:** Security logs must be protected from modification.

Lipner's Full Matrix Integrity Model

Category	ISP (System Program)	IO (Operational)	ISL (System Low)
ID (Development)	SSD (System Development)	SD (Development)	SD (Development)
IP (Production)	SP (Production)	SP (Production)	SD (Development)

Key Points:

- **Integrity Levels:**
 - ID (Development) → Lower integrity (less trusted).
 - IP (Production) → Higher integrity (more trusted).

- Access Control:
 - System Development (SSD) → **Modifies system programs.**
 - Development (SD) → Can access but not modify production.
 - Production (SP) → Can operate and modify production.
- Ensures:
 - Separation of duties.
 - Protection of production data.
 - Controlled development and deployment.

Example Scenario:

- A developer writes software in Development (D).
- The code is reviewed & approved, then moved to Production Code (PC).
- A user runs the program, which accesses Production Data (PD), but cannot modify it.
- The Audit Manager monitors logs for any violations.

6. Clark-Wilson Integrity Model

- Key Concepts:
 - CDIs (Constrained Data Items): Data with integrity controls.
 - UDIs (Unconstrained Data Items): Data without integrity controls.
 - IVPs (Integrity Verification Procedures): Ensure CDIs are valid.
 - TPs (Transaction Procedures): Move system from one valid state to another.
- Certification Rules (CR):
 - CR1: IVPs must ensure all CDIs are valid.
 - CR2: TPs must transform CDIs from one valid state to another.
 - CR3: Enforces separation of duties.
 - CR4: TPs must log operations.
 - CR5: TPs must validate UDIs before converting them into CDIs.
- Enforcement Rules (ER):
 - ER1: Only certified TPs can modify CDIs.
 - ER2: Users must be associated with specific TPs and CDIs.
 - ER3: Users must be authenticated before executing TPs.
 - ER4: Certifiers of TPs cannot execute them, enforcing separation of duty.
- Comparison with Biba Model:
 - Biba: Multilevel integrity, no explicit certification.

- Clark-Wilson: Focuses on transactions, separation of duties, and certified transformations.

Ensures data integrity via well-defined processes, transaction validation, and separation of duties.

HYBRID POLICIES : Overview of Security Models

Chinese Wall Model

- Goal: Prevent **conflicts of interest** in business consulting and financial sectors.
- Solution: Group companies into **Conflict of Interest (COI) classes** and control access.
- Example:
 - Tony advises Bank A; he cannot advise Bank B if they are in the same COI class.

Rules:

1. CW-Simple Security Rule (Read Rule)
 - A user can read only within a company dataset (CD) they already accessed OR
 - They must not have accessed any competing company's dataset.
2. **CW-Property (Write Rule)*
 - A user can write only if all the datasets they read belong to the same company.
3. Sanitization:
 - Public (sanitized) data is freely accessible.

Comparison to Other Models

- Bell-LaPadula: CW has no security labels, but tracks past access history.
- Clark-Wilson: If a person uses different processes, CW rules may be bypassed.

Clinical Information Security (CISS Policy)

- Focus: **Integrity & confidentiality** in medical records.

Key Principles:

1. Access Control List (ACL): Only approved clinicians can read/write records.
2. Responsible Clinician: Can add others to ACL and must notify patients.
3. Audit Logs: Track who accessed what, when, and why.
4. Data Deletion: Records cannot be deleted until the appropriate time.
5. **Data Sharing (Confinement Principle)**: Information can be copied only if the new record's ACL is a subset of the original.

6. Aggregation Protection: **Patients must be notified** if someone gains access to many records (prevents misuse).
7. Enforcement & Auditing: Must be independently evaluated.

Comparison to Other Models

- Bell-LaPadula: Similar **confinement rules**, but B-LP focuses on subjects, while CISS focuses on objects (medical records).
- Clark-Wilson:
 - Medical records = CDIs (Constrained Data Items).
 - Updates = TPs (Transformation Procedures).
 - Integrity checks ensure only clinicians validate records.
 - Auditing ensures records are append-only.

□ Chinese Wall Model prevents **conflicts of interest** in business.

□ CISS Policy ensures medical record security through access control and auditing.

□ RBAC & ORCON **provide flexible access control mechanisms**.

□ Bell-LaPadula vs. Clark-Wilson:

- B-LP controls who can access data.
- C-W controls how **data is modified and ensures** integrity.

ORCON (Originator Controlled Access Control)

- Goal: Control how information is shared and copied **after creation**.
- Example: **A government official writes a memo and must approve any further sharing.**

Rules:

1. Originator (X) controls dissemination.
2. No sharing with new organizations without originator's approval.
3. Copies inherit the same access restrictions.

Why MAC & DAC Fail?

- DAC issue: The owner can freely change permissions, violating rule #2.
- MAC issue: Requires creating new categories for every access scenario, making it impractical.

Solution: Hybrid Approach

- MAC: Access restrictions cannot be changed by the owner.
- DAC: Originator can grant or deny permissions per subject/object.

RBAC (Role-Based Access Control)

- Goal: Access depends on job role, not individual identity.
- Example:
 - Allison (Math Dept bookkeeper) can access **financial records**.
 - **When Betty replaces her, Betty gets access, not Allison.**

Key Concepts:

- Role (r): A set of job functions.
- Transactions (trans(r)): Actions allowed for a role.
- Active Role (actr(s)): Role a subject is currently using.
- Authorized Roles (authr(s)): Roles a subject can take.

Rules:

1. Role Assignment: A subject must have a role to execute a transaction.
2. Role Authorization: A subject can only assume an assigned role.
3. Transaction Authorization: **A subject can only perform actions** allowed for its role.

Advanced Features:

- Role Containment: Higher roles include permissions of lower roles (Trainer > Trainee).
- Separation of Duty: Some roles must be mutually exclusive (e.g., a person cannot be both an auditor and an accountant).
- ORCON: Controls who can share copies of information (Hybrid of MAC & DAC).
- RBAC: Controls **who can perform which actions** based on job roles.
- Separation of Duty: **Prevents conflicts by restricting multiple roles for a single user.**

NON-INTERFERENCE and POLICY COMPOSITION

Composition Problem in security policies:

- Composition problem: How to merge multiple security policies into one coherent and secure policy.
- **Important for distributed systems, merged organizations, and integrated networks.**
- **Challenges arise when policies conflict or interact unexpectedly.**

1. Problem Statement

- Two organizations with different security policies merge.
- How do we create a unified, consistent security policy?
- Can the new policy maintain security guarantees of individual policies?

2. Policy Composition Challenges

Noninterference

- **Definition:** A high-security level should not affect a low-security level's outputs.
- **Example:** A government system should prevent classified info (HIGH) from influencing unclassified outputs (LOW).

HIGH Inputs Affect LOW Outputs

- **Risk:** If high-level actions cause changes in observable low-level behavior, a security breach occurs.
- **Example:** A CPU scheduling difference due to a classified process may leak information to an unclassified observer.

Nondeducibility

- **Definition:** A LOW-level observer should not be able to infer HIGH-level inputs from observable outputs.
- **Example:** If a stock market system delays updates based on insider trades (HIGH), external traders (LOW) might infer secret transactions.

Restrictiveness

- When can policies be successfully merged?
 1. **Autonomy Principle:** If access is allowed in any original policy, it should be allowed in the merged policy.
 2. **Security Principle:** If access is forbidden in any original policy, it must be forbidden in the merged policy.

3. Policy Composition Process

Composition of Bell-LaPadula

- **Why?** Some security standards require secure components to form a secure system.
- **Key Question:** Under what conditions is this secure?

Assumptions

- Systems implement security policies precisely.
- **Must compose security lattices to merge labels and access rules.**

Lattice Composition Example

- Assume clearances: $LOW < S < HIGH < TS$.
- Assume categories: SOUTH, EAST, WEST.

- Merged system must reflect relationships among different levels and categories.

Different Policies

- What does "secure" mean in the merged system?
- Which component's policy takes priority?
- Solutions:
 1. **Autonomy** → Allow access if any original policy allowed it.
 2. **Security** → Forbid access if any original policy forbade it.

4. Policy Composition Example (Gong & Qian)

- **System X**: Bob cannot access Alice's files.
- **System Y**: Eve and Lilith can access each other's files.
- **Merged Policy**: Bob can access Eve's files, Lilith can access Alice's files.

Key Question: Can Bob access Lilith's files?

Solution (Gong & Qian Approach)

1. Define Access Sets (AS):

- $AS(X) = \emptyset$
- $AS(Y) = \{(Eve, Lilith), (Lilith, Eve)\}$
- $AS(X \cup Y) = \{(Bob, Eve), (Lilith, Alice), (Eve, Lilith), (Lilith, Eve)\}$

2. Compute Transitive Closure (All Implied Accesses):

- $AS(X \cup Y)^+ = \{(Bob, Eve), (Bob, Lilith), (Bob, Alice), (Eve, Lilith), (Eve, Alice), (Lilith, Eve), (Lilith, Alice)\}$

3. Delete Conflicting Accesses: Remove (Bob, Alice) since it conflicts with System X.

4. Result: Bob can access Lilith's files.

- Composition can **introduce unintended accesses**.
- Transitive closure **must be computed**.
- Forbidden accesses must be **explicitly removed**.

5. Interference & Covert Channels

Example: Covert Channel via CPU Usage

- Scenario:
 - Holly (HIGH) & Lara (LOW) use separate virtual machines.
 - They share the same CPU, but load-based scheduling creates an **observable pattern**.
 - Holly runs a program to create high CPU usage → Lara detects it → Secret information leaks.

- Solution Considerations:
 - Prevent covert channels by strict resource allocation.
 - Modify policy definitions to include unintended interference.

6. Formal Model for Security Composition

- System as State Machine:
 - Subjects $S = \{s_i\}$
 - States $\Sigma = \{\sigma_i\}$
 - Commands $Z = \{z_i\}$
 - Outputs $O = \{o_i\}$
- Functions:
 - State Transition Function: $T: C \times \Sigma \rightarrow \Sigma$ (Defines system evolution)
 - Output Function: $P: C \times \Sigma \rightarrow O$ (Defines what is visible)
- Projection & Purge Mechanism:
 - Used to remove unauthorized outputs when merging policies.

7. Key Insights

- Composition is Hard:
 - Combining security policies often allows unintended behaviors.
 - Some policies may dominate others, affecting security guarantees.
- Mathematical Modeling is Key: Using state machines, projections, and closures helps reason about composition security.

NON INTERFERENCE :

Noninterference

- intuition: outputs visible to one user should only depend on inputs they can see
- formal definition: if a group G executes commands in set A , they are noninterfering with another group G' if outputs for G' remain unchanged
- notation: $A, G :| G'$

Example:

- initial state $\sigma_0 = (0,1)$, commands = [(Heidi, xor0), (Lucy, xor1), (Heidi, xor1)]
- checking projection for Lucy shows Heidi's commands affect her outputs, so noninterference is false
- modifying Heidi's commands to affect only H-bit ensures noninterference holds

Security policy

- defines authorized and unauthorized states
- ensures no forbidden interferences occur
- can be expressed as a set of noninterference assertions

Output-consistency

- states are output-consistent if, for a given command, they produce the same output for subjects in its domain
- ensures that executing a command does not leak information to unauthorized domains

Protection domains

- system partitions into protection domains $D = \{d_1, d_2, \dots, d_n\}$
- relation r defines information flow rules between domains

Projection function

- π' removes commands that do not interfere with a domain
- if executing a command interferes with domain d , it remains visible in $\pi'd$

Noninterference security

- a system is noninterference-secure if executing commands produces the same output as their projection
- ensures no unauthorized information flow

Lemma and proof

- if state transitions remain equivalent after projection and output-consistency holds, system is noninterference-secure
- follows from the definition of output-consistency and noninterference

UNWINDING THEOREM :

- Links security of sequences of state transitions to security of individual transitions
- Helps prove ML security by checking if system meets conditions from known lemmata
- Does not guarantee system security due to implementation and operational issues locally respects
- system X locally respects policy r if applying a command c does not affect domain d when it should not
- ensures no unintended side effects on unrelated domains

Transition-consistent

- system is transition-consistent if applying command c does not change the equivalence of states under r
- ensures consistency in how commands affect the system

Lemma

- if two commands c_1 and c_2 allow information flow into d , their execution order does not matter
- $T^*(c_1c_2, \sigma) = T(c_1, T(c_2, \sigma)) = T(c_2, T(c_1, \sigma))$

Theorem

- if a system is output-consistent, transition-consistent, and locally respects r , then it is **noninterference-secure** under r
- provides a basis for verifying system designs against noninterference policies

Proof outline

1. use induction on the sequence of commands cs
2. base case: cs is empty, holds trivially
3. assume hypothesis holds for cs
4. for cs followed by $cn+1$, consider two cases:
 - $\text{dom}(cn+1)rd$ holds: apply transition consistency and induction hypothesis
 - $\text{dom}(cn+1)rd$ does not hold: apply local respect property
5. conclude that $T^*(cs, \sigma_0) \sim_d T^*(\pi'd(cs), \sigma_0)$
6. since X is output-consistent, it follows that X is noninterference-secure

Access Control Matrix

Example of Interpretation

- Given: Access control information
- Question: Are given conditions enough to provide noninterference security?
- Assume: System in a particular state
- Encapsulates: Values in ACM

ACM Model

- Objects: $L = \{ l_1, \dots, l_m \}$ (Locations in memory)
- Values: $V = \{ v_1, \dots, v_n \}$ (Values that L can assume)
- Set of States: $\Sigma = \{ \sigma_1, \dots, \sigma_k \}$
- Set of Protection Domains: $D = \{ d_1, \dots, d_j \}$

Functions

- $\text{value}: L \times \Sigma \rightarrow V \rightarrow$ Returns value v stored in location l when system in state σ
- $\text{read}: D \rightarrow 2^L \rightarrow$ Returns set of objects observable from domain d
- $\text{write}: D \rightarrow 2^L \rightarrow$ Returns set of objects writable from domain d

Interpretation of ACM

- Functions represent ACM
- Subject s in domain $d = \text{dom}(c)$, object o
 - $r \in A[s, o] \Leftrightarrow o \in \text{read}(d)$
 - $w \in A[s, o] \Leftrightarrow o \in \text{write}(d)$

Equivalence Relation

$\sigma_a \sim_{\text{dom}(c)} \sigma_b \Leftrightarrow \forall l_i \in \text{read}(d)$
 $\text{value}(l_i, \sigma_a) = \text{value}(l_i, \sigma_b)$

- You can read exactly the same locations in both states, where $\sigma_b = T(c, \sigma_a)$

Enforcing Policy r

Five Requirements

- Three general ones: Describe dependence of commands on rights over input and output
 - Hold for all ACMs and policies
- Two specific ones: Apply to some security policies
 - Hold for most policies

First Requirement

- Output of command c executed in $\text{dom}(c)$ depends only on values for which subjects in $\text{dom}(c)$ have read access
- $\sigma_a \sim_{\text{dom}(c)} \sigma_b \Rightarrow P(c, \sigma_a) = P(c, \sigma_b)$

Second Requirement

- If c changes l_i , then c can only use values of objects in $\text{read}(\text{dom}(c))$ to determine the new value
- If l_i can't be read in $\text{dom}(c)$, then values in l_i may differ after c is applied to σ_a or σ_b

Third Requirement

- If c changes l_i , then $\text{dom}(c)$ provides subject executing c with write access to l_i
- $\text{value}(l_i, T(c, \sigma_a)) \neq \text{value}(l_i, \sigma_a) \Rightarrow l_i \in \text{write}(\text{dom}(c))$

Fourth Requirement

- If domain u can interfere with domain v , then every object that can be read in u can also be read in v
- $u \text{ r } v \Rightarrow \text{read}(u) \subseteq \text{read}(v)$

Fifth Requirement

- If subject s can write object o in v and subject s' can read o in u , then domain v can interfere with u
- $l_i \in \text{read}(u) \text{ and } l_i \in \text{write}(v) \Rightarrow v \text{ r } u$

Theorem

- If X satisfies the five conditions, then X is noninterference-secure with respect to r
- Proof: Must show X is:
 1. Output-consistent
 2. Locally respects r
 3. Transition-consistent
- Then, by the Unwinding Theorem, the theorem holds

Output-Consistency

- Take equivalence relation \sim_d , first condition is the definition of output-consistent

Local Respect for r

- Proof by contradiction
 - Assume $(\text{dom}(c), d) \notin r$ but $\sigma_a \sim_d T(c, \sigma_a)$ does not hold
 - Some object has value changed by c
 - Condition 3: $l_i \in \text{write}(d)$
 - Condition 5: $\text{dom}(c) \text{ r } d$, contradiction
 - So, $\sigma_a \sim_d T(c, \sigma_a)$ holds, meaning X locally respects r

Transition-Consistency

- Assume $\sigma_a \sim_d \sigma_b$
- Must show $\text{value}(l_i, T(c, \sigma_a)) = \text{value}(l_i, T(c, \sigma_b))$ for $l_i \in \text{read}(d)$

Case 1

- $\text{value}(l_i, T(c, \sigma_a)) \neq \text{value}(l_i, \sigma_a)$
- Condition 3: $l_i \in \text{write}(\text{dom}(c))$
- Condition 5: $\text{dom}(c) \text{ r } d$
- Condition 4: $\text{read}(\text{dom}(c)) \subseteq \text{read}(d)$
- Apply Condition 2, so $T(c, \sigma_a) \sim_{\text{dom}(c)} T(c, \sigma_b)$

Case 2

- $\text{value}(l_i, T(c, \sigma_b)) \neq \text{value}(l_i, \sigma_b)$
- Follows the same steps as Case 1

Case 3

- Neither of the previous two cases
- $\text{value}(l_i, T(c, \sigma_a)) = \text{value}(l_i, \sigma_a)$
- $\text{value}(l_i, T(c, \sigma_b)) = \text{value}(l_i, \sigma_b)$
- Then $T(c, \sigma_a) \sim_d T(c, \sigma_b)$

Conclusion: In all three cases, X is transition-consistent

Policies Changing Over Time

- Problem: Previous analysis assumes a static system
- Reality: ACM changes as system commands are issued
- Solution: **Condition noninterference on $\text{cando}(w, s, z)$**
- Example: If $\neg \text{cando}(w, \text{Lara}, \text{"write f"})$, then Lara can't interfere with any user by writing f

Generalized Noninterference

- $G \subseteq S$ (Group of subjects)
- $A \subseteq Z$ (Set of commands)
- p (Predicate over elements of C^*)
- π'' : Removes commands from A executed by G when p holds

Intuition

- $\pi''(cs) = cs$, except when p holds
- If p holds, commands from A by members of G are replaced with v
- **Effect**: Similar to deleting entries from cs

Noninterference Definition

- Users in G executing A are noninterfering with G' under condition p iff:
 - $\forall cs \in C^*, \forall s \in G', \text{proj}(s, cs, \sigma_i) = \text{proj}(s, \pi''(cs), \sigma_i)$
- Notation: $A, G :| G'$ if p

Example: Security Policy

- If s can't execute z, then s can't use z to interfere
- $\forall (s \in S) \forall (z \in Z) [\{z\}, \{s\} :| S \text{ if } \neg \text{cando}(w, s, z)]$

Noninterference with Right Transfers

- Problem: Rights can be passed dynamically
- Example: *pass(s, z)* gives *s* right to execute *z*
- Policy:
 - No subject *s* can use *z* to interfere if, in the previous state, *s* did not have the right to *z*, and no subject gave it to *s*
- Effect:
 - Ensures that a subject's first execution of *z* does not affect other subjects' observation of the system

Key Ideas in Transition Consistency and Noninterference

Transition Consistency

- If two states, σ_a and σ_b , are indistinguishable under a dependency *d*, their transitions should also be indistinguishable.
- This is shown by checking three cases where a transition $T(c, \sigma)$ affects a variable *li*:
 1. If *li* changes → Conditions ensure σ_a and σ_b remain consistent.
 2. If *li* changes in the other state → Similar reasoning applies.
 3. If *li* doesn't change → The property holds naturally.

Policies Changing Over Time

- Real systems update access controls dynamically, unlike static models.
- Example: A user *w* gets new permissions, affecting their ability to execute commands.
- Noninterference Rule: If a user can't execute a command, they can't interfere with others.

Noninterference & Security Policies

- Generalized Noninterference: Users in *G* executing commands in *A* don't affect users in *G'*, ensuring security.
- Example: If a user lacks write access, their actions shouldn't impact other users.
- Extending to Dynamic Rights: If a user gains permission at time *t*, it must be verified how they got it.

Security and Composition

- Combining secure systems can break security if timing issues arise.
- Example: A low-level observer (Lara) might infer high-level actions (Heidi) from indirect effects.

- Feedback-Free Systems: If a system has no loops, combining secure systems remains secure.

Nondeducibility vs. Strong Noninterference

- Nondeducibility: Low-level users can't infer high-level inputs.
- **Strong Noninterference**: Even the timing of high-level events can't be observed.
- Example: If a system leaks information via subtle timing differences, it fails strong noninterference.

Conclusion

- Security models ensure users can't infer hidden data.
- Dynamic policies and composition of systems introduce challenges.
- Solutions: Feedback-free structures, delayed updates, and strict access rules maintain security.

Restrictiveness & Composition in Security Systems

Restrictiveness: A system is restrictive if it satisfies four key properties (previously discussed).

Composition

- Idea: If two restrictive systems are combined, the result should also be restrictive.
- Why? By properties 3 and 4, a high-level output followed by a low-level output behaves like a low-level input, so security should hold.

Composite System

- System M1's outputs become M2's inputs.
- A composite system consists of paired states from M1 and M2:
 - $\mu_{1i}, \mu_{2i} \rightarrow$ States of M1, M2.
 - $(\mu_{1i}, \mu_{2i}) \rightarrow$ A state in the composite system.

Transition Conditions

An event e causes a transition in the composite system if any of these hold:

1. M1 changes: M1 moves from $\mu_{1a} \rightarrow \mu_{1b}$, e doesn't affect M2 ($\mu_{2a} = \mu_{2b}$).
2. M2 changes: M2 moves from $\mu_{2a} \rightarrow \mu_{2b}$, e doesn't affect M1 ($\mu_{1a} = \mu_{1b}$).
3. Both change: M1 moves $\mu_{1a} \rightarrow \mu_{1b}$ and M2 moves $\mu_{2a} \rightarrow \mu_{2b}$, where e is an input to one and an output from the other.

Intuition

- Any event in the composite system must trigger at least one component's transition.

- If only one component changes, then the event must be irrelevant to the other component when they are not connected.

Equivalence for Composite Systems

- Define an equivalence relation:
 - $(\sigma a, \sigma b) \equiv_C (\sigma c, \sigma d)$ if $\sigma a \equiv \sigma c$ and $\sigma b \equiv \sigma d$.
- This matches the equivalence relation in Property 2 for individual components.

Conclusion

- Restrictive systems remain restrictive when composed properly.
- Transitions follow predictable rules, ensuring secure information flow.