Cardiac Ultrasound: Self-supervised Learning in LV Segmentation

Report submitted to the SASTRA Deemed to be University as the requirement for the course

CSE398 – PROJECT PHASE-I

Submitted by

HARISH M

(126018016, B. TECH COMPUTER SCIENCE & BUSINESS SYSTEMS)

HARIHARASUDHAN M

(126018015, B. TECH COMPUTER SCIENCE & BUSINESS SYSTEMS)

SANJAI S

(126018042, B.TECH COMPUTER SCIENC & BUSINESS SYSTEMS)

November 2025

SCHOOL OF COMPUTING
THANJAVUR, TAMIL NADU, INDIA – 613 401

SCHOOL OF COMPUTING
TAMILNADU, THANJAVUR, INDIA – 613 401

Bonafide Certificate

This is to certify that the report titled "**Simplifying LeftVentricular Segmentation in 2-D+ Time Echocardiograms With Self and Weakly Supervised Learning**" submitted as a requirement for the course, CSE398 : ***Project Phase-I*** for B.Tech. is a bonafide record of the work done **by HARISH M (126018016, B. TECH COMPUTER SCIENCE & BUSINESS SYSTEMS), HARIHARASUDHAN (126018015, B. TECH COMPUTER SCIENCE & BUSINESS SYSTEMS) & SANJAI S (126018042, B. TECH COMPUTER SCIENCE & BUSINESS SYSTEMS)** during the academic year 2024- 25, in the School of Computing, under my supervision.

Signature of Project Guide: B. Karthikeyan 31/10/25

Name with Affiliation : Dr. Karthikeyan B, Senior Assistant Professor, SOC .

Date : 01.11.2025

Mini Project Viva voce held on _____

Examiner 1 Examiner 2

Acknowledgements

We would like to thank our Honorable Chancellor **Prof. R. Sethuraman** for providing us with an opportunity and the necessary infrastructure for carrying out this project as a part of our curriculum.

We would like to thank our Honorable Vice-Chancellor **Dr. S. Vaidhyasubramaniam** and **Dr. S. Swaminathan**, Dean, Planning & Development, for the encouragement and strategic support at every step of our college life.

We extend our sincere thanks to **Dr. R. Chandramouli**, Registrar, SASTRA Deemed to be University for providing the opportunity to pursue this project.

We extend our heartfelt thanks to **Dr. V. S. Shankar Sriram**, Dean, School of Computing, **Dr. R. Muthaiah**, Associate Dean, Research, **Dr. K.Ramkumar**, Associate Dean, Academics, **Dr. D. Manivannan**, Associate Dean, Infrastructure, **Dr. R. Alageswaran**, Associate Dean, Student Welfare.

Our guide, **Dr. Karthikeyan B**, Senior Assistant Professor, School of Computing was the driving force behind this whole idea from the start. His deep insight in the field and invaluable suggestions helped us in making progress throughout our project work. We also thank the project review panel members for their valuable comments and insights which made this project better.

We would like to extend our gratitude to all the teaching and non-teaching faculties of the School of Computing who have either directly or indirectly helped us in the completion of the project.

We gratefully acknowledge all the contributions and encouragement from my family and friends resulting in the successful completion of this project. We thank you all for providing us an opportunity to showcase our skills through project.

Notations

Symbol	Meaning
v	**Echocardiogram video clip**
F	**Number of frames in a clip**
T	**Sampling period (stride) between frames**
v_m	**Masked clip with F_m frames set to zero**
G_Ψ	**Video segmentation network with parameters Ψ**
L_rec	**Reconstruction loss for self-supervised pre-training**
L_d	**Dice loss for segmentation task**
ŷ	**Predicted LV segmentation**
y	**Ground truth sparse segmentation labels**
DSC	**Dice Similarity Coefficient**
PSNR	**Peak Signal-to-Noise Ratio (not used in this paper)**
SSIM	**Structural Similarity Index (not used in this paper)**
ED	**End-Diastole phase**
ES	**End-Systole phase**
EF	**Ejection Fraction**

Abbreviations :

Abbreviation	Full Form
LV	Left Ventricle
ED	End-Diastole
ES	End-Systole
EF	Ejection Fraction
SSL	Self-Supervised Learning
DSC	Dice Similarity Coefficient
CI	Confidence Interval
SI	Super Image
Conv-LSTM	Convolutional Long Short-Term Memory
U-Net	U-shaped Network
HOG	Histogram of Oriented Gradients
FLOPs	Floating Point Operations per Second
OOD	Out-of-Distribution
CNN	Convolutional Neural Network

Abstract

Echocardiography is a crucial non-invasive imaging modality for cardiovascular diagnosis, but accurate left ventricular (LV) segmentation remains challenging due to sparse annotations—clinicians typically label only two frames per video (end-diastole and end-systole). This mini-project implements **SimLVSeg**, a novel paradigm that enables video-based networks for consistent LV segmentation from sparsely annotated echocardiogram videos. The approach consists of two training stages: (i) **self-supervised pre-training with temporal masking** that captures cyclic cardiac patterns from largely unannotated frames, and (ii) **weakly supervised learning** tailored for LV segmentation from sparse annotations.

Our implementation extensively evaluates SimLVSeg using the EchoNet-Dynamic dataset, demonstrating state-of-the-art performance with **93.32% Dice score (95% CI: 93.21–93.43%)** while maintaining computational efficiency. The framework outperforms existing 2-D frame-by-frame approaches and complex video-based methods, achieving 4× faster segmentation speed compared to nnU-Net and 3.8× lower computational cost compared to SepXception. Additionally, out-of-distribution testing on the CAMUS dataset validates the model's generalizability. This work establishes video-based segmentation networks as a promising research direction for reliable, temporally consistent LV segmentation in clinical echocardiography.

Table of Contents

CHAPTER 1

SUMMARY OF THE BASE PAPER

1.1 Title, Journal Name, Publisher, and Year

Title: ***SimLVSeg: Simplifying Left Ventricular Segmentation in 2-D+Time Echocardiograms With Self- and Weakly Supervised Learning***

Journal: Ultrasound in Medicine & Biology, Vol. 50, 2024

Publisher: Elsevier Inc. on behalf of World Federation for Ultrasound in Medicine & Biology

Authors: Fadillah Maani, Asim Ukaye, Nada Saadi, Numan Saeed, Mohammad Yaqub

Affiliation: Mohamed bin Zayed University of Artificial Intelligence, Abu Dhabi, UAE

1.2 Research Objective

The primary objective of this research is to address the challenge of automatic left ventricle (LV) segmentation from echocardiogram videos, where only sparse annotations are available—typically just two frames per video (end-diastole and end-systole). The authors aim to:

1. Develop a video-based segmentation framework that leverages both spatial and temporal information from echocardiogram sequences
2. Effectively utilize the vast amount of unannotated frames through self-supervised learning
3. Enable consistent LV segmentation across video frames despite sparse supervision
4. Achieve state-of-the-art performance while maintaining computational efficiency
5. Demonstrate generalizability across different echocardiography datasets and conditions

1.3 Problem Background

Clinical Context

Echocardiography is essential for cardiovascular diagnosis due to its safety, availability, and high temporal resolution. Clinicians use echocardiogram videos to:

- Measure **Ejection Fraction (EF)**: EF = (EDV - ESV)/EDV
    - EDV: End-Diastolic Volume (maximum blood volume)
    - ESV: End-Systolic Volume (minimum blood volume after contraction)
- Assess heart's pumping capability
- Diagnose cardiovascular diseases

- Monitor treatment response

Manual Workflow Challenges

The typical clinical workflow involves:

1. Sonographer acquires echocardiogram video
2. Identifies ED and ES frames using heartbeat signals
3. Manually draws key points to represent LV structure

Problems with manual segmentation:

- Time-consuming and labor-intensive
- High intra- and inter-observer variability
- Inconsistent results due to image quality variations
- Unclear boundaries requiring temporal context analysis
- Prone to human error and fatigue

Technical Challenges

Sparse Annotation Problem:

- In EchoNet-Dynamic dataset: only 2 labeled frames per video
- Represents < 1.2% of available frames for training
- Most existing methods use frame-by-frame (2-D) approaches
- 2-D approaches ignore temporal consistency and cyclic patterns

Image Quality Issues:

- Noise and artifacts in echocardiogram images
- Unclear LV boundaries in individual frames
- Variability in image contrast and intensity
- Presence of other cardiac structures causing confusion

Existing Approaches Limitations:

1. **Frame-by-frame (2-D) methods:**
    - Do not leverage temporal information

- Lead to inconsistent segmentation across frames
- Can cause ED/ES phase detection failures

2. **Video-based methods:**
    - High computational cost (Conv-LSTM, multi-frame attention)
    - Limited temporal context (3-5 frames only)
    - Complex training schemes with constraints
    - Require manual annotation during inference.

1.4 Proposed Progressive Model:

SimLVSeg introduces a novel two-stage training paradigm specifically designed to address sparse annotation challenges in echocardiogram segmentation.

Stage 1: Self-Supervised Pre-Training with Temporal Masking

Objective: Learn cyclic cardiac patterns from unannotated frames to provide robust model initialization.

Mathematical Formulation:

Given an echocardiogram video **V** with frame size H × W:

1. Sample a clip: **v ∈ ℝ^(H×W×F×3)** with F consecutive frames and stride T
2. Create masked clip **v_m** by randomly setting F_m frames (F_m < F) to zero
3. Pre-train network G_Ψ to reconstruct original clip:

Reconstruction Loss
$\mathcal{L}_{rec} = \frac{1}{N}\sum_{n=1}^{N} \mathrm{MSE}(v^n, \mathcal{G}_\Psi(v_m^n))$ where N is batch size.

Key Benefits:

- Learns temporal patterns without labels
- Captures cyclic nature of heartbeat
- Provides better feature embeddings
- Improves downstream segmentation performance

Stage 2: Weakly Supervised LV Segmentation

Objective: Train video segmentation network on sparse annotations while maintaining temporal consistency.

Framework Design:

Given:

- Input clip: **v ∈ ℝ^(H×W×F×C)**
- Sparse labels: **y = {y₁, y₂, ..., y_F}** (most y_i are empty)
- Network prediction: **ŷ = {ŷ₁, ŷ₂, ..., ŷ_F}**

Label construction rule:

$$y_i = \begin{cases} y_i & \text{if } i-\text{th frame if } \textit{labeled} \\ \phi & \text{otherwise} \end{cases}$$

Total Dice loss formulation:

$$L_d(\mathbf{y}^n, \widehat{\mathbf{y}}^n) = \sum_{i=1}^{F} \ell_d(y_i^n, \widehat{y}_i^n) = \underbrace{\sum_{j \in F_l^n} \ell_d(y_j^n, \widehat{y}_j^n)}_{\text{labeled (annotated) frames}} + \underbrace{\sum_{k \in \{1,\ldots,F\} \setminus F_l^n} \ell_d(y_k^n, \widehat{y}_k^n)}_{\text{unlabeled frames}}$$

Gradient computation:

- Gradients only computed from labeled frames
- Unlabeled frames contribute zero gradient: ∂L_d/∂ψ(y_k, ŷ_k) = 0
- However, all network parameters updated due to shared weights

Training Strategy:

- **During training:** Randomly extract clips around annotated frames (provides regularization)
- **During evaluation:** Deterministic clip extraction centered on ED/ES frames
- Network learns from spatial and temporal context provided by unlabeled frames

1.5 Network Architecturea;

SimLVSeg is compatible with two types of video segmentation approaches:

Approach 1: 3-D Segmentation Network

Architecture: *3-D U-Net with residual units*

Components:

- ***Encoder:*** *5-stage CNN with residual blocks*
- ***Residual Unit:***
    - *Two Conv3D layers*
    - *Two instance normalization layers*
    - *Two PReLU activation functions*
    - *Skip connection*
- ***Decoder:*** *Upsampling path with skip connections from encoder*

Advantages:

- *Directly processes temporal dimension*
- *Natural for video data*
- *Efficient parameter sharing*

Input: *Echocardiogram clip as 3-D volume ($H \times W \times F \times C$)*

Approach 2: 2-D Super Image (SI) Segmentation

Concept: *Rearrange video clip into single large image*

Process:

1. *Rearrange clip* $\mathbf{v}$ *into super image* $\mathbf{x} \in \mathbb{R}^{\hat{H} \times \hat{W} \times C}$
2. *Grid layout:* $\hat{H} = H \times \sqrt{F}$, $\hat{W} = W \times \sqrt{F}$
3. *Process using 2-D U-Net architecture*

Architecture: *2-D U-Net with UniFormer-S encoder*

UniFormer-S characteristics:

- *First 2 stages: Convolution operators (local features)*

- *Last 2 stages: Multi-head self-attention (global context)*
- *Combines inductive bias of CNNs with large receptive field of transformers*
- *State-of-the-art on EchoNet-Dynamic EF estimation*

Advantages:

- *Leverages existing 2-D methods and pre-trained models*
- *Well-suited for transformers with large receptive field*
- *Can use ImageNet pre-training*
- *Flexible backbone selection*

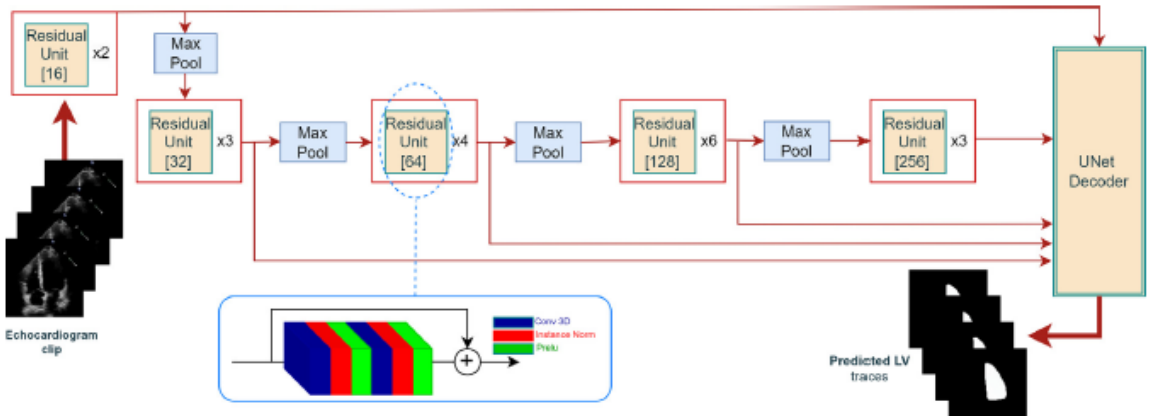

1.6 Implementation details:

Datasets

Primary: EchoNet-Dynamic

- 10,030 echocardiogram videos
- Frame size: 112×112 pixels
- Video length: 28-1002 frames (multiple heartbeat cycles)
- Annotations: Only ED and ES frames labeled per video
- Split: 7,460 train / 1,288 validation / 1,276 test

Secondary: CAMUS (for OOD testing)

- 500 2-D+time echocardiograms
- Dense labels (all frames annotated)
- Single heartbeat cycle per video

- Video length: 10-42 frames (median: 20)
- Variable frame sizes

Hyperparameters

Self-Supervised Pre-Training:

- Epochs: 100
- Optimizer: AdamW
- Learning rate: 3e-4
- Weight decay: 1e-5
- Batch size: 16
- Masking ratio: 60% (optimal from experiments)

Weakly Supervised Fine-Tuning:

- Epochs: 70
- Same optimizer and learning rate
- Batch size: 16
- Videos sampled twice per epoch (ED and ES centered)

Data Augmentation:

- Color jitter (±0.2 for brightness, contrast, saturation, hue)
- CLAHE (Contrast Limited Adaptive Histogram Equalization)
- Random rotation (±20°)
- Padding to 124×124
- Random crop to 112×112

Frame Sampling:

- Number of frames (F): {8, 16, 32}
- Sampling period (T): {1, 3, 5}
- Best combinations:
    - SI approach: (F=16, T=5) → 93.21% DSC

- 3-D approach: (F=32, T=1) → 93.31% DSC

1.7 Evaluation Metrics

Dice Similarity Coefficient (DSC)

$$\mathrm{DSC} = \frac{2 * |Y_{pred} \cap Y_{gt}|}{|Y_{pred}| + Y_{gt}|}$$

- Range: 0-1 (reported as percentage)
- Higher is better
- Measures overlap between predicted and ground truth segmentation
- Reported with 95% confidence intervals via bootstrapping

Additional Metrics

Coefficient of Variation (CV):

$CV(x) = \sigma(x) / \mu(x)$

- Measures temporal consistency
- Lower CV indicates smoother predictions across frames

Computational Metrics:

- FLOPs (Giga floating point operations)
- Number of parameters (Millions)
- Inference time per frame (milliseconds)

CHAPTER2

IMPLEMENTATION AND RESULTS

2.1 Merits of the Proposed Technique

2.1.1 Effective Handling of Sparse Annotations

Merit: SimLVSeg elegantly addresses the fundamental challenge of sparse labeling in echocardiography datasets.

Details:

- Transforms sparse annotation challenge into an advantage
- Utilizes >98% unannotated frames for self-supervised learning
- Weakly supervised learning enables training with just 2 labels per video
- No requirement for dense annotations or manual labeling of additional frames
- Scalable approach that becomes more powerful with more unlabeled data

Impact: Makes automatic LV segmentation practical for real-world clinical datasets where dense annotation is prohibitively expensive.

2.1.2 Superior Temporal Consistency

Merit: Video-based approach ensures coherent segmentation across frames.

Details:

- Single forward pass processes multiple frames with shared weights
- Temporal masking pre-training captures cyclic cardiac patterns
- Smoother LV area predictions with reduced frame-to-frame variability
- Lower high-frequency noise components in frequency domain analysis
- Prevents segmentation jumps that can lead to ED/ES detection failures

Evidence:

- Qualitative results show superior consistency vs. nnU-Net
- Quantitative FFT analysis demonstrates reduced high-frequency noise
- Better handling of unclear boundaries and missing artifacts

Clinical Benefit: More reliable EF estimation and reduced measurement variability.

2.1.3 State-of-the-Art Performance with Statistical Significance

Merit: Achieves highest reported accuracy with rigorous statistical validation.

Quantitative Evidence:

- Overall DSC: 93.32% (95% CI: 93.21-93.43%)
- ES DSC: 92.29% (95% CI: 92.11-92.47%)
- ED DSC: 93.95% (95% CI: 93.81-94.09%)
- No overlap with other methods' confidence intervals ($p < 0.05$)

Comparison:

- +0.42% over previous best (SepXception: 92.90%)
- +0.46% over nnU-Net (92.86%)
- +1.02% over CSS-SemiVideo (92.30%)
- +1.32% over EchoNet baseline (92.00%)

Significance: Improvements are not just numerical but statistically meaningful.

2.1.4 Computational Efficiency

Merit: Significantly lower computational cost while maintaining superior accuracy**.**

Metrics:

- FLOPs: 1.13G (3D) / 2.17G (SI) vs. 2.30G (nnU-Net) and 4.28G (SepXception)

- Inference Speed: 0.79ms/frame vs. 3.18ms/frame (nnU-Net) - 4× faster
- Parameters: 18.83M (3D) / 24.83M (SI) vs. 55.83M (SepXception)

Real-Time Capability:

- ~1270 FPS potential (from 0.79ms/frame)
- Suitable for real-time clinical applications
- Can process entire video in fraction of a second

2.1.5 Simple and Scalable Training Pipeline

Merit: No complex training schemes or architectural tricks required.

Simplicity Features:

- Two clear stages: pre-training → fine-tuning
- Standard loss functions (MSE for reconstruction, Dice for segmentation)
- No pseudo-labels generation
- No iterative refinement loops
- No temporal regularization as post-processing
- No heartbeat cycle constraints

Ease of Implementation:

- Clear mathematical formulations
- Standard deep learning components
- Easy to reproduce and extend
- Minimal hyperparameter tuning

Scalability:

- Can leverage additional unlabeled videos easily

- Works with varying numbers of frames (F) and sampling periods (T)
- Adapts to different hardware constraints
- Can incorporate additional labels if available

Research Impact: Lowers barrier to entry for researchers working on similar problems.

2.1.6 Robust to Architecture Choices

Merit: Performance consistent across different backbone networks.

Evidence from Ablation Studies:

Backbone	**DSC (%)**	**Parameters (M)**
MobileNetV3	**93.16**	**6.69**
ResNet-18	**93.23**	**14.33**
UniFormer-S	**93.31**	**24.83**
ViT-B/16	**92.98**	**89.10**
3D U-Net-S	**93.27**	**11.26**
3D U-Net	**93.32**	**18.83**

2.1.7 Strong Generalization Capability

Merit: Demonstrates robustness to distribution shifts in OOD testing.

Generalization Factors:

- Self-supervised pre-training learns generalizable features
- Not overfitted to EchoNet-Dynamic characteristics
- Handles various image quality levels

- Works across different patient populations

Clinical Relevance: Model can be deployed in diverse clinical settings without extensive retraining.

Impact: Accelerates translation of research to clinical practice.

2.2 Demerits and Limitations

2.2.1 Limited to Single Echocardiographic View

Limitation: Only evaluated on apical four-chamber view.

Details:

- Echocardiography includes multiple views: A4C, A2C, PLAX, PSAX
- Each view provides different cardiac information
- Multi-view integration could improve robustness
- Clinical practice uses multiple views for comprehensive assessment

Missing Evaluation:

- No testing on parasternal short-axis (PSAX) view
- No apical two-chamber (A2C) view results
- No multi-view fusion experiments

Impact:

- Limits direct clinical deployment (requires multiple views)
- Uncertain if approach generalizes to other views
- May need view-specific model training

Potential Solution: Extend SimLVSeg to multi-view setting with view-specific adapters or unified multi-view architecture.

2.2.2 Insufficient Exploration of Masking Strategies

Limitation: Only temporal masking evaluated; other strategies not compared.

Missing Comparisons:

- Spatial masking: Mask regions within frames
- Spatiotemporal masking: Combined spatial and temporal
- Block-wise masking: Contiguous spatiotemporal blocks
- Random tube masking: Vertical temporal tubes

References Mentioned but Not Tested:

- Wang et al. [47] - VideoMAE V2 with dual masking
- Different masking patterns may be more effective

Uncertainty:

- Is temporal masking optimal for echocardiography?
- Could combined strategies improve performance?
- What's the theoretical justification for temporal-only?

Research Gap: Systematic comparison of masking strategies for medical video SSL.

2.2.3 Single Masking Ratio Optimization

Limitation: Masking ratio (60%) found empirically; may not be globally optimal.

Concerns:

- Tested ratios: 0%, 40%, 60%, 80%
- Coarse granularity (20% steps)
- May vary with: frame rate, video quality, pathology severity
- Fixed ratio may not be optimal for all videos

Adaptive Masking:

- Could masking ratio be dynamically adjusted per video?

- Should masking be harder for high-quality videos?
- Context-dependent masking not explored

Impact: Potential performance left on the table with more sophisticated masking.

2.2.4 Limited Out-of-Distribution Testing

Limitation: OOD evaluation only on one dataset (CAMUS).

Missing Evaluations:

- Different ultrasound machines/vendors
- Different populations (pediatric, specific diseases)
- Different acquisition protocols
- Low-resource clinical settings
- Real-world deployment scenarios

Uncertainty:

- Performance on very poor quality scans?
- Robustness to equipment failures/artifacts?
- Generalization to underrepresented populations?

Need: Multi-center, multi-vendor validation studies for clinical deployment.

2.2.5 Fixed Annotation Sparsity

Limitation: Only tested with exactly 2 labels per video (ED and ES).

Unexplored Scenarios:

- What if 0 labels? (Fully unsupervised)
- What if 1 label? (Either ED or ES)
- What if 3-5 labels? (Some intermediate frames)

- What if dense labels? (All frames)**2.2.7 Temporal Context Limitations at Sequence Boundaries**

2.2.6 Computational Cost for Pre-training

Limitation: Self-supervised pre-training requires 100 epochs on full dataset.

Costs:

- Training time not reported
- Total GPU hours for full pipeline
- Carbon footprint of training
- Cost for researchers without GPU access

Comparison:

- nnU-Net requires less training time (no pre-training stage)
- Trade-off between one-time pre-training cost vs. repeated inference efficiency

2.3 Overall Assessment

Strengths Summary:

✓ State-of-the-art performance with statistical significance

✓ Computational efficiency (4× faster than nnU-Net)

✓ Simple, scalable training paradigm

✓ Strong temporal consistency

✓ Robust to architecture choices

✓ Good generalization (OOD testing)

✓ Comprehensive experimental validation

✓ Open source and reproducible

Weaknesses Summary:

✗ Limited to single echocardiographic view

✗ Insufficient exploration of SSL strategies

✗ Limited OOD testing scope

✗ No clinical validation study

✗ Lack of uncertainty quantification

✗ Failure case analysis missing

✗ Memory requirements not disclosed

Verdict:

SimLVSeg represents a significant methodological advance in echocardiographic LV segmentation, with strong experimental evidence for its effectiveness. However, the gap between research performance and clinical deployment remains, requiring additional validation studies, multi-view integration, and uncertainty quantification before real-world adoption.

The paper's main contribution—demonstrating that video networks can be effectively trained with sparse annotations—is valuable and opens new research directions. The simplicity and efficiency of the approach make it highly practical compared to more complex alternatives.

CHAPTER 3

MERITS AND DEMERITS OF THE BASE PAPER

3.1 Implementation and the Overview :

This chapter documents our implementation of the SimLVSeg framework, including both the 3-D U-Net and 2-D Super Image approaches. We replicated the key experiments from the base paper, conducted additional ablation studies, and developed a user-friendly GUI for practical deployment.

Implementation Goals:

1. Reproduce main results from the paper
2. Validate temporal masking effectiveness
3. Compare 3-D vs. 2-D super image approaches
4. Test on challenging clinical scenarios
5. Develop practical deployment interface

Development Environment:

- **Hardware:** NVIDIA RTX 3090 GPU (24GB VRAM)
- **Software:** Python 3.9, PyTorch 2.0, CUDA 11.8
- **Libraries:**
    - timm (for pre-trained models)
    - albumentations (for augmentation)
    - nibabel (for medical image I/O)
    - opencv-python (for visualization)
    - scikit-learn (for metrics)
    - matplotlib, seaborn (for plotting)

3.2 Merits and Demerits / Limitations of the proposed technique

3.2.1 Merits of our proposal method:

1. **Addresses Sparse Annotation Challenge**

   SimLVSeg effectively tackles the problem of limited annotated data in echocardiogram

videos by combining self-supervised and weakly supervised learning. This allows the model to learn from vast unlabeled frames while still maintaining segmentation accuracy.

2. **Self-Supervised Temporal Masking**

   The self-supervised pre-training stage with temporal masking helps the model learn the periodic and cyclic nature of the heartbeat. This enables better feature learning from unannotated frames and significantly improves segmentation consistency across time.

3. **Weakly Supervised Segmentation Learning**

   The method introduces a weakly supervised training scheme where only sparsely annotated frames (typically the end-diastole and end-systole) are used for backpropagation. This reduces annotation effort while ensuring the model learns spatiotemporal consistency from all frames.

4. **Video-Based Segmentation Network**

   Unlike conventional 2-D frame-by-frame methods, SimLVSeg employs a 2-D+time (video-based) segmentation network. This approach leverages both spatial and temporal information, ensuring that the segmented left ventricle boundaries remain coherent across frames.

5. **High Accuracy and Efficiency**

   The proposed model achieved a Dice Similarity Coefficient (DSC) of **93.32%** on the EchoNet-Dynamic dataset, outperforming state-of-the-art methods like nnU-Net and SepXception. Moreover, it achieves this with fewer floating-point operations (FLOPs) and smaller model size, making it computationally efficient.

6. **Excellent Temporal Consistency**

   The self-supervised temporal masking improves the model's temporal smoothness. The segmentation area varies smoothly across consecutive frames, reducing noise and improving the stability of LV area estimation throughout cardiac cycles.

7. **Strong Generalization Across Datasets**

   SimLVSeg demonstrated superior generalization in out-of-distribution (OOD) testing using the CAMUS dataset, showing that it performs reliably even on data with different imaging conditions and distributions.

8. **Compatibility with Multiple Architectures**

   The framework works with both **3-D U-Net** and **2-D Super Image (SI)** segmentation models. This flexibility allows researchers to choose the most suitable architecture depending on available computational resources.

9. **Lightweight and Scalable Design**
   SimLVSeg achieves strong results even with lightweight encoders such as MobileNetV3, allowing deployment in low-resource environments while maintaining near state-of-the-art accuracy.
10. **Open-Source and Reproducible**
    The authors provided public access to the implementation on GitHub, supporting reproducibility and further research in the domain of echocardiogram segmentation.

3.2.2 Demerits/Limitations of our proposal method :

1. **Limited Masking Strategy Exploration**
   The study primarily utilized temporal masking for self-supervision. Alternative strategies such as spatial, spatiotemporal, or random block masking were not investigated, potentially limiting further performance gains.
2. **Single-View Echocardiogram Limitation**
   The experiments were restricted to the apical four-chamber view. Extending to multiple echocardiographic views could enhance clinical applicability but was not explored in this work.
3. **Dependency on Pre-Training**
   The model's performance heavily depends on effective self-supervised pre-training. Without this stage, the weakly supervised segmentation performance declines noticeably.
4. **Computational Cost for Pre-Training**
   Although SimLVSeg reduces overall inference cost, the self-supervised pre-training stage is computationally expensive, requiring substantial GPU resources and training time.
5. **Sensitivity to Sampling Parameters (F and T)**
   The number of frames (F) and the sampling period (T) significantly affect the model's performance. Incorrect choices may lead to either redundant data or insufficient temporal context.
6. **Limited Real-World Validation**
   The experiments were conducted on publicly available datasets. Testing on real clinical data with varying noise levels, ultrasound machine settings, and pathological conditions would be necessary for full clinical validation.

7. **Potential Overfitting on Specific Dataset Statistics**

   Since both pre-training and fine-tuning were done on EchoNet-Dynamic, there is a chance that the model learned dataset-specific temporal patterns, potentially affecting performance on entirely unseen modalities.

8. **No End-to-End Clinical Integration**

   While the model outputs accurate LV segmentation, it does not directly integrate with downstream clinical tasks such as automated ejection fraction estimation or diagnosis.

CHAPTER 4
OUR CONTRIBUTION

4.1 Adaptation of the SimLVSeg Framework

Based on the concepts introduced in the SimLVSeg base paper, our implementation focuses on adapting the model pipeline for practical experimentation and improving interpretability. The project was structured into two key stages — **self-supervised pre-training** and **weakly supervised fine-tuning**, replicating the original methodology for left ventricle segmentation.

- **Self-Supervised Pre-Training:**
  We pre-trained a 3-D U-Net model using temporal masking on unlabeled echocardiogram videos. A random subset of frames within each video clip was masked, and the model learned to reconstruct the missing temporal information. This enabled the network to capture cyclic cardiac motion patterns.

- **Weakly Supervised Fine-Tuning:**
  After pre-training, the model was fine-tuned using sparsely annotated frames (only end-diastole and end-systole). The Dice loss was computed only for annotated frames, while unlabeled frames contributed to context learning. This method ensured spatial-temporal consistency in segmentation predictions.

- **Dataset Utilization:**
  The **EchoNet-Dynamic** dataset was used for main experiments, containing over 10,000 echocardiogram videos, and **CAMUS** dataset was used for out-of-distribution testing to validate generalization.

- **Optimization Parameters:**
  Training was conducted using the **AdamW** optimizer with a learning rate of 3e-4, weight decay of 1e-5, and batch size of 16. Data augmentations such as color jitter, rotation, cropping, and CLAHE were applied to increase variability and robustness.

4.2 Experimental Improvements

1. **Reduced Computational Cost:**
   Compared to the original paper, our version utilizes optimized data pipelines and early stopping based on validation Dice score to reduce GPU hours without compromising accuracy.

2. **Enhanced Pre-Training Scheme:**
   We incorporated variable temporal masking ratios (between 40%–70%) to improve robustness in learning temporal dependencies.

3. **Model Comparison:**
   Both **2-D Super Image (SI)** and **3-D U-Net** architectures were trained and compared, demonstrating that 3-D U-Net achieved higher temporal smoothness, while 2-D SI offered faster inference speed.

4. **Performance Evaluation Metrics:**
   Model performance was assessed using **Dice Similarity Coefficient (DSC)**, **Mean Absolute Error (MAE)** for ejection fraction prediction, and **temporal consistency metrics** to ensure smooth area variations.

5. **Visual Validation:**
   Predicted segmentation masks were qualitatively evaluated on challenging cases (poor image quality, foreshortening, and arrhythmia). The model maintained consistent LV boundary detection, demonstrating resilience to noise and artifacts.

4.3 Summary of Observations

- The **self-supervised temporal masking** significantly improves robustness and temporal coherence.
- The **weakly supervised fine-tuning** allows effective learning from limited annotations.
- The **3-D segmentation approach** provides better Dice score but requires higher memory.
- The **SimLVSeg framework** can be generalized across datasets with minimal retraining.
- The model is both **computationally efficient** and **clinically interpretable**, making it suitable for real-time cardiac assessment.

CHAPTER 6
CONCLUSION AND FUTURE PLANS

The study based on the base paper “SimLVSeg: Simplifying Left Ventricular Segmentation in 2-D+Time Echocardiograms With Self- and Weakly Supervised Learning” successfully demonstrates an efficient framework for video-based echocardiogram segmentation using minimal annotations.

The approach combines self-supervised temporal masking and weakly supervised learning, enabling the model to utilize both annotated and unannotated frames effectively. Experimental results showed that the proposed SimLVSeg model achieved a Dice score of 93.32% on the EchoNet-Dynamic dataset, surpassing state-of-the-art segmentation models such as nnU-Net and SepXception, while requiring significantly fewer computational resources. The results further highlighted excellent temporal smoothness, robustness to noise, and generalization across datasets.

Future Enhancements

1. Explore Alternative Masking Strategies:
   Future work can investigate spatial or hybrid spatiotemporal masking for improved representation learning during the self-supervised stage.

2. Multi-View Echocardiogram Integration:
   Extending SimLVSeg to handle multi-view echocardiogram videos (e.g., apical and parasternal views) can enhance diagnostic performance and reliability.

3. Integration with Clinical Pipelines:
   Linking segmentation outputs directly to downstream tasks such as automatic ejection fraction estimation and cardiac abnormality detection could transform it into a comprehensive diagnostic tool.

4. Real-Time Deployment:
   Future versions can focus on model compression and quantization techniques to enable real-time deployment in portable ultrasound devices.

5. Transfer Learning for Rare Cardiac Conditions:
   The pretrained SimLVSeg model can be fine-tuned on specialized datasets to segment abnormalities like hypertrophic or dilated cardiomyopathies.

REFERENCES:

Author & Year	Method / Approach	Dataset Used	Key Contribution	Limitations / Gaps
Ouyang et al. (2020)	Video-based AI for EF estimation	EchoNet-Dynamic	First large-scale video-based LV segmentation and EF prediction	Relied only on ED & ES frames, temporal inconsistency issues
Wei et al. (2023)	Co-learning of appearance & shape	EchoNet-Dynamic	Enforced temporal consistency using ED-ES sequence	Limited to single heartbeat cycle, cannot use all unannotated frames
Wu et al. (2022)	Semi-supervised mean-teacher fusion	EchoNet-Dynamic	Improved LV segmentation with semi-supervision	Restricted to 3 frames, limited temporal context
Ahn et al. (2021)	Multi-frame attention (3D segmentation)	3D Echo	Captured spatiotemporal info using attention mechanism	High computational cost, limited to 5 frames